

Big Data Analytics

Session 5(b) **Cross Validation**

So far

- Compute **MSE/error rate on the training data**
 - Easy!
- Calculate **MSE/error rate on the test data**
 - Easy, if the designated test set is available
 - ➔ Unfortunately, this is usually not the case
- **Training MSE/error rate** can **dramatically underestimate** the test MSE/error rate.
- **Main question**: How to estimate the test MSE/error rate in the absence of the designated test data?

Cross Validation

- Solution:
 - Estimate the test error rate by
 - holding out a subset of the training observations from the fitting process, and then
 - applying the statistical learning method to those held out observations.

Training data for fitting the model

Training data
for fitting

← Held out data for testing

Outline




- Cross Validation on **Regression Problems**
 1. The Validation Set Approach
 2. Leave-One-Out Cross Validation
 3. K-fold Cross Validation
 - Bias-Variance Trade-off for k-fold Cross Validation
- Cross Validation on **Classification Problems**

1. The Validation Set Approach



- Suppose that we would like **estimate the test error** associated with fitting a particular statistical learning method
- We can achieve this goal by **randomly splitting** the data into
 - **training part** and
 - **validation (testing, or hold-out) part**



Example: Auto Data




- Suppose that we want to predict **mpg** from **horsepower**
 - Linear model:
 - $\text{mpg} \sim \text{horsepower}$
 - How to do it?
 - Randomly split **Auto** data set (392 obs.) into training (196 obs.) and validation data (196 obs.)

```
> set.seed(1)
> train=sample(392,196)
```
 - Fit the model using the training data set

```
> lm.fit.train=lm(mpg~horsepower, data=Auto, subset=train)
```
 - Then, evaluate the model using the validation data set

```
> mean((Auto$mpg-predict(lm.fit.train, Auto))[-train]^2)
[1] 26.14142
```
- Plot the observations and linear relationship between mpg and horsepower

Did you get this?



A way to improve



- From the plot, there appears to be a **non-linear relationship** between **mpg** and **horsepower**.
- Try the quadratic model: $\text{mpg} \sim \text{horsepower} + \text{horsepower}^2$
- Repeat the procedure
 - Randomly split **Auto** data set (392 obs.) into training (196 obs.) and validation data (196 obs.) – **the same as before**
 - Fit the model using the training data set

```
> lm.fit2.train=lm(mpg~poly(horsepower,2),data=Auto, subset=train)
```
 - Then, evaluate the model using the validation data set

```
> mean((Auto$mpg-predict(lm.fit2.train,Auto))[-train]^2)
[1] 19.82259                                #linear model: 26.14142
```
- Compare the two test errors
 - The quadratic model has a smaller test error, thus is better!

Results: Auto Data

- Left: Validation error rate for a single split
- Right: Validation method repeated 10 times, each time the split is done randomly!
- There is a lot of variability among the MSE's... Not good! We need more stable methods!



The Validation Set Approach



- Advantages:

- Simple
- Easy to implement

- Disadvantages:

- The validation MSE can be highly variable
- Only a subset of observations are used to fit the model (training data).


Statistical methods tend to perform worse when trained on fewer observations.

2. Leave-One-Out Cross Validation (LOOCV)



- This method is **similar** to the Validation Set Approach, but it tries to **address the latter's disadvantages**.
- For each suggested model, do:
 - Split the data set of size n into
 - Training data set (blue) size: $n - 1$
 - Validation data set (beige) size: 1
 - Fit the model using the training data
 - Validate model using the validation data, and compute the corresponding MSE
 - Repeat this process n times
 - The MSE for the model is computed as follows:

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n MSE_i.$$



LOOCV vs. Validation Set Approach



- LOOCV has **less bias**
 - We repeatedly fit the statistical learning method using training data that contains $n - 1$ obs., i.e. **almost all the data set is used**
- LOOCV produces a **less variable MSE**
 - The **validation set approach** produces different MSE when applied repeatedly due to randomness in the splitting process
 - Performing **LOOCV** multiple times will always yield the same results, because we split based on 1 obs. each time
- LOOCV is **computationally intensive** (disadvantage)
 - We fit a model n times!

Perform LOOCV in R



- Using the Auto data set again, building a linear model

```
> glm.fit=glm(mpg~horsepower,data=Auto)
```

```
># This is the same as lm.fit(mpg~horsepower,data=Auto)
```

```
> library(boot) #cv.glm() is in the boot library
```

```
> cv.err=cv.glm(Auto,glm.fit)
```

```
> # cv.glm() does the LOOCV
```

```
> cv.err$delta
```

```
[1] 24.23151 24.23114
```


The MSE is 24.23151.

3. k-fold Cross Validation

- LOOCV is computationally intensive, so we can run k -fold Cross Validation instead
- With k -fold CV, we divide the data set into k different parts (e.g. $k = 5$, or $k = 10$, etc.)
- We then remove the first part, fit the model on the remaining $k-1$ parts, and see how good the predictions are on the left out part (i.e. compute the MSE on the first part)
- We then repeat this k different times taking out a different part each time
- By averaging the k different MSE's we get an estimated validation (test) error rate for new observations

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k \text{MSE}_i.$$

K-fold Cross Validation



Perform K-fold CV in R



- Very easy!

```
> glm.fit=glm(mpg~horsepower,data=Auto)
```

```
># This is the same as in LOOCV
```

```
> library(boot) # This is the same as in LOOCV
```

```
> cv.err=cv.glm(Auto,glm.fit, K=10)
```

```
#K means K-fold, can be 5, 10 or other numbers
```



```
> cv.err$delta
```

```
[1] 24.3120 24.2926
```

The MSE is 24.3120.



Auto Data: LOOCV vs. k-fold CV

- Left: LOOCV error curve
- Right: 10-fold CV was run many times, and the figure shows the slightly different CV error rates
- LOOCV is a special case of k -fold, where $k = n$
- They are both stable, but LOOCV is more computationally intensive!





Auto Data: Validation Set Approach vs. k-fold CV Approach

- Left: Validation Set Approach
- Right: 10-fold Cross Validation Approach
- Indeed, 10-fold CV is more stable!



K-fold Cross Validation on the Simulated Data



- Blue: True Test MSE
 - Black: LOOCV MSE
 - Orange: 10-fold MSE
 - Refer to chapter 2 for Fig 2.9. More example see Fig 5.6
- Model: Smoothing spline

Bias-Variance Trade-off for k -fold CV



- Putting aside that LOOCV is more computationally intensive than k -fold CV... Which is better LOOCV or k -fold CV?
 - LOOCV is **less bias** than k -fold CV (when $k < n$)
 - LOOCV: uses $n-1$ observations
 - K -fold CV: uses $(k-1)n/k$ observations
 - But, LOOCV has **higher variance** than k -fold CV (when $k < n$)
 - The mean of many highly correlated quantities has higher variance
 - Thus, there is a **trade-off** between what to use
- Conclusion:
 - We tend to use k -fold CV with ($k = 5$ and $k = 10$)
 - These are the magical k 's ☺
 - It has been empirically shown that they yield test error rate estimates that suffer **neither from excessively high bias**, **nor from very high variance**

Cross Validation on Classification Problems



- So far, we have been dealing with CV on regression problems
- We can use cross validation in a classification situation in a similar manner
 - Divide data into k parts
 - Hold out one part, fit using the remaining data and compute the **error rate** on the hold out data
 - Repeat k times
 - CV error rate is the average over the k errors we have computed

LAB

The Validation Set Approach



- Goal: To estimate the test MSE
- Approach:
 - Randomly pick half of the data as the train data
 - Estimate the test MSE on the other half

- Random (but repeatable): `set.seed(1)`
- Pick half: `train = sample(392, 196)`

- Test MSE:
- Other half: set indices as `[-train]`

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

- y_i : mpg
- $\hat{f}(x_i)$:
 - `lm.fit=lm(mpg~horsepower, data=Auto, subset=train)`
 - `Predict(lm.fit, Auto)`
- `1/n: mean()`
- `mean((mpg-predict(lm.fit, Auto))[-train]^2)`

The Validation Set Approach



- Linear regression (Degree 1)
 - `lm.fit = lm(mpg~horsepower, data=Auto, subset=train)`
 - `mean((mpg-predict(lm.fit, Auto))[-train]^2)`
 - 26.14
- Polynomial regression
 - Degree 2
 - `lm.fit2 = lm(mpg~poly(horsepower, 2), data=Auto, subset=train)`
 - `mean((mpg-predict(lm.fit2, Auto))[-train]^2)`
 - 19.82
 - Degree 3
 - `lm.fit3 = lm(mpg~poly(horsepower, 3), data=Auto, subset=train)`
 - `mean((mpg-predict(lm.fit3, Auto))[-train]^2)`
 - 19.78
- What can we conclude from the above results?

The Validation Set Approach



- Choosing a different training set, then we will obtain different errors on the validation set.

```
set.seed(2) or set.seed(i) (i≠1)
```

```
repeat the rest
```

- Notice the variability on the results

Leave-One-Out Cross Validation



- Function `glm()`

- In logistic regression:

```
glm(y~x, family="binomial", data=..)
```

- In linear regression:

```
glm(y~x, data=..)
```

```
> glm.fit=glm(mpg~horsepower, data=Auto)
> coef(glm.fit)
(Intercept)    horsepower
      39.936         -0.158
```

the same as

```
lm(y~x, data=..)
```

```
> lm.fit=lm(mpg~horsepower, data=Auto)
> coef(lm.fit)
(Intercept)    horsepower
      39.936         -0.158
```

Leave-One-Out Cross Validation



- Function `cv.glm()` in `boot` library
 - Produces a list with several components, including the cross-validation estimate for the test error: `delta`
 - `cv.glm(data, glmfit, cost, K)`

```
> library(boot)
> glm.fit=glm(mpg~horsepower, data=Auto)
> cv.err=cv.glm(Auto, glm.fit)
> cv.err$delta
      1      1
24.23 24.23
```

- `Delta` is a vector of length two. For LOOCV, the two are the same.

Leave-One-Out Cross Validation



- Experiment on the CV for increasingly complex polynomial fits
- Initialise a vector of length `len` all to be number `val`
 - `vec = rep(val, len)`
- Using for loop to repeat procedure

```
> cv.error=rep(0,5)
> for (i in 1:5){
+   glm.fit=glm(mpg~poly(horsepower,i),data=Auto)
+   cv.error[i]=cv.glm(Auto,glm.fit)$delta[1]
+ }
> cv.error
[1] 24.23 19.25 19.33 19.42 19.03
```

K-Fold Cross Validation

- Implement k-fold CV by passing the argument K

```
> set.seed(17)
> cv.error.10=rep(0,10)
> for (i in 1:10){
+   glm.fit=glm(mpg~poly(horsepower,i),data=Auto)
+   cv.error.10[i]=cv.glm(Auto,glm.fit, K=10)$delta[1]
+ }
> cv.error.10
[1] 24.21 19.19 19.31 19.34 18.88 19.02 18.90 19.71 18.95 19.50
```

- The two numbers associated with delta
 - The first number is the raw/standard CV estimate of prediction error
 - The second number is the adjusted CV estimate. The adjustment is designed to compensate for the bias introduced by not using leave-one-out cross-validation