

Big Data Analytics

Session

Data Stream Mining

Outline



- Introduction
- Data Stream Classification
- Data Stream Clustering
- Novel Class Detection


Outline




- Introduction
- Data Stream Classification
- Data Stream Clustering
- Novel Class Detection

Introduction

- Characteristics of Data streams are:
 - Continuous flow of data
 - Examples



Network traffic



Sensor data



Call center records

Challenges




- Infinite length
- Concept-drift
- Concept-evolution
- Feature evolution


Infinite length

- Impractical to store and use all historical data
 - Requires infinite storage

- And running time



Concept-Drift




Negative instance •

Positive instance ○

Instances victim of concept-drift ●

Concept-Evolution

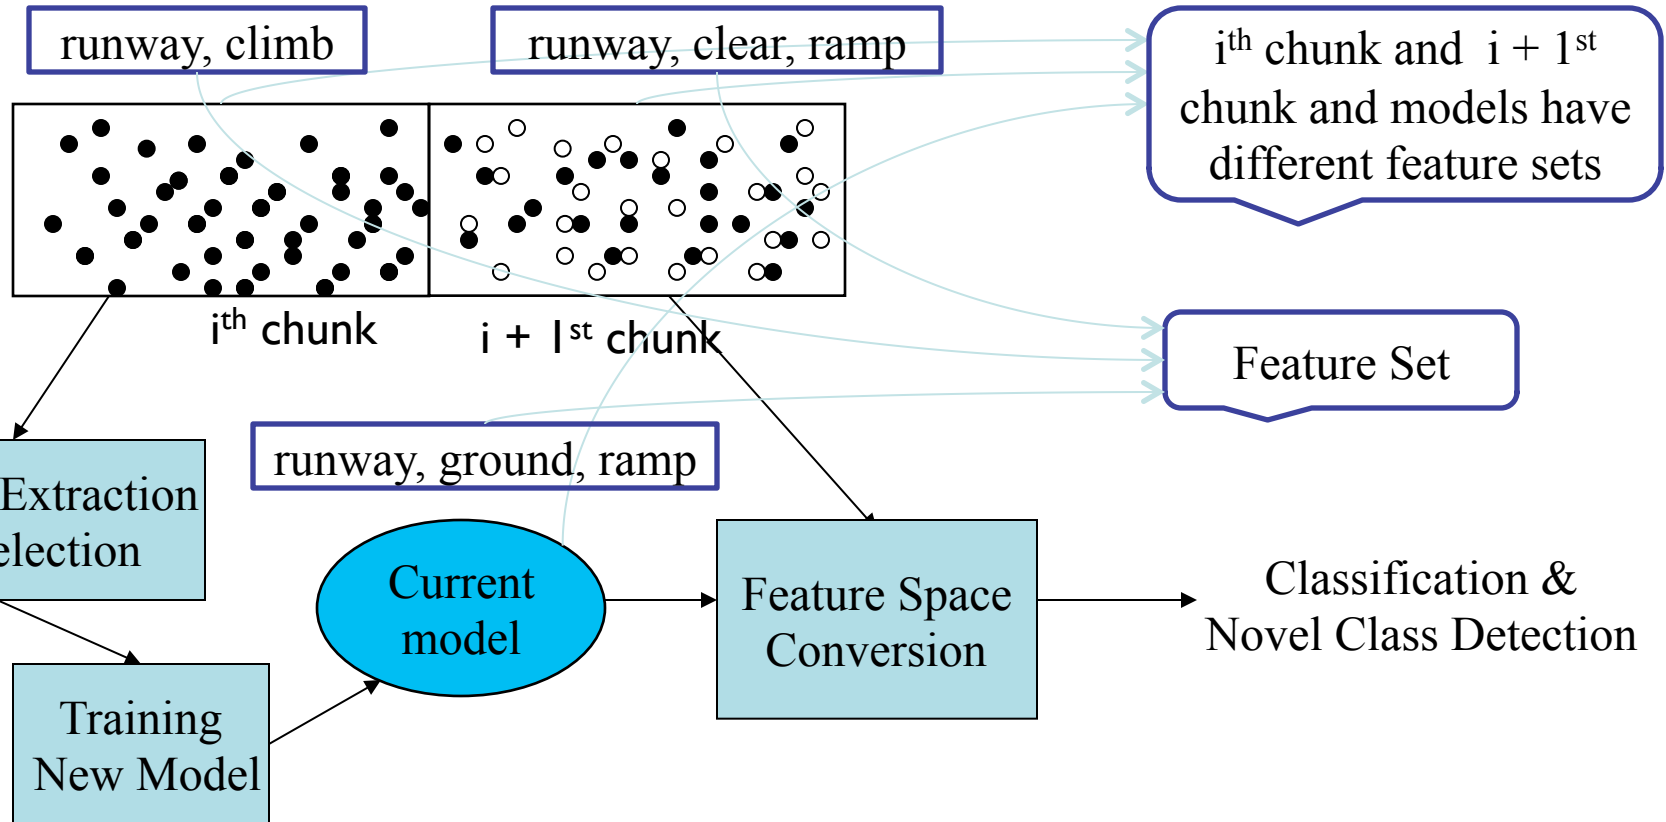


- Classification rules:
 - R1. if $(x > x_1 \text{ and } y < y_2)$ or $(x < x_1 \text{ and } y < y_1)$ then class = +
 - R2. if $(x > x_1 \text{ and } y > y_2)$ or $(x < x_1 \text{ and } y > y_1)$ then class = -
- Existing classification models misclassify novel class instances

Dynamic Features

- Why new features evolving
 - Infinite data stream
 - Normally, global feature set is unknown
 - New features may appear
 - Concept drift
 - As concept drifting, new features may appear
 - Concept evolution
 - New type of class normally holds new set of features
- Different chunks may have different feature sets

Dynamic Features



- Existing classification models need complete fixed features and apply to all the chunks.
- Global features are difficult to predict.
- One solution is using all English words and generate vector.
- Dimension of the vector will be too high.

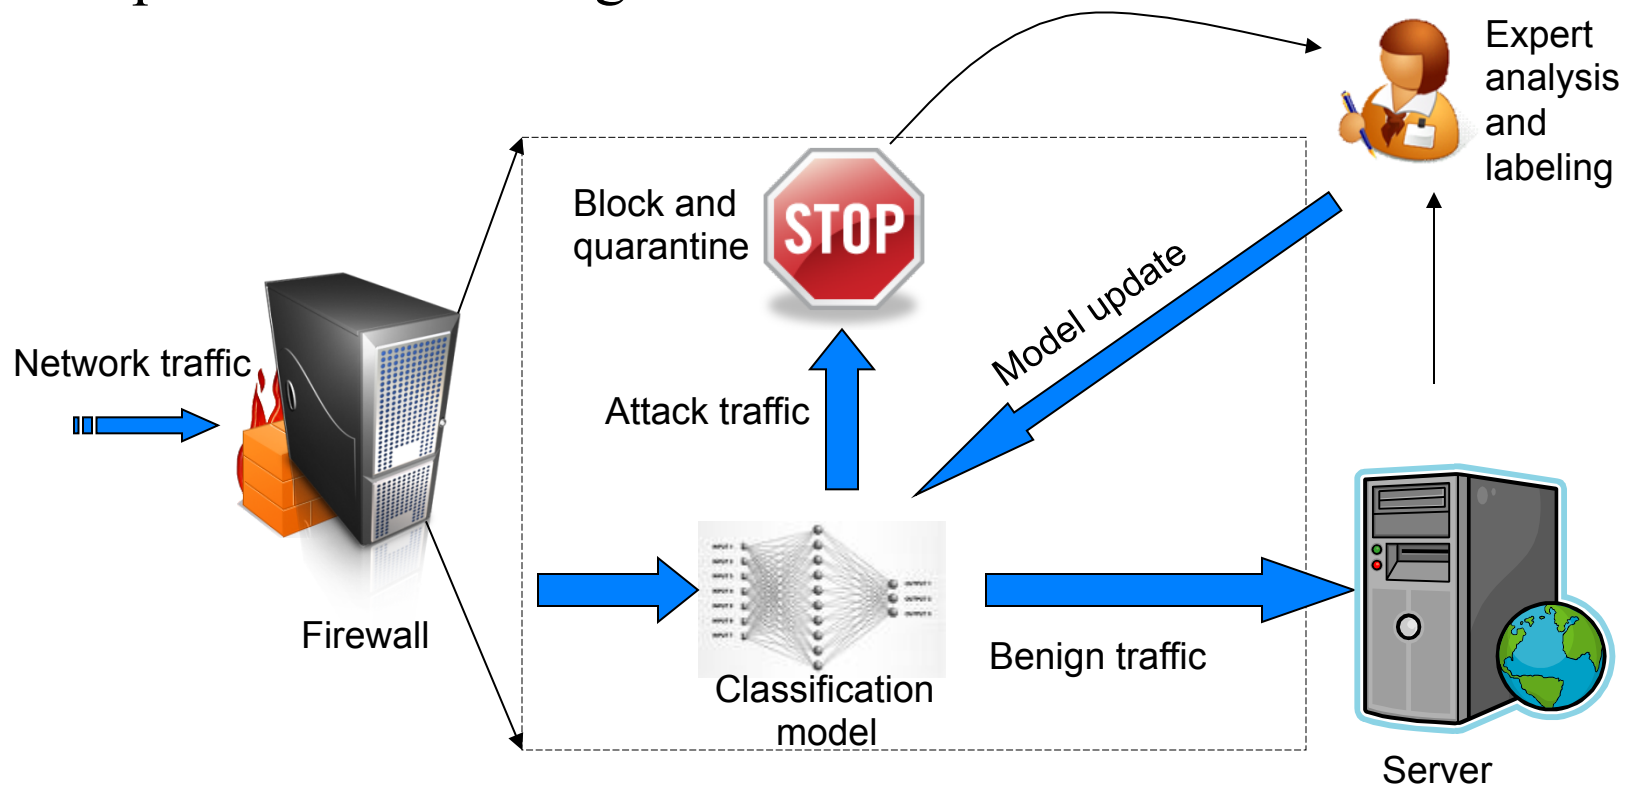
Outline



- Introduction
- Data Stream Classification
- Data Stream Clustering
- Novel Class Detection

Data Stream Classification

- Uses past *labeled data* to build classification model
- Predicts the labels of future instances using the model
- Helps decision making



Data Stream Classification – Applications



- Security monitoring
- Network monitoring and traffic engineering
- Business: credit card transaction flows
- Telecommunication calling records
- Web logs and web page click streams
- Financial market: stock exchange

Data Stream Classification – Approaches



- Single model incremental classification
- Ensemble – Model based classification
 - Supervised
 - Semi-supervised
 - Active learning

Data Stream Classification – Approaches



- Single model incremental classification
- Ensemble – Model based classification
 - Supervised
 - Semi-supervised
 - Active learning

Single Model Incremental Classification



- Introduction
- Problems in mining Data Stream
- Decision Tree
- Very Fast Decision Tree (VFDT)
- Concept-Adapting VFDT (CVFDT)

Problems in Mining Data Streams



- Traditional data mining techniques usually require entire data set to be present.
- Random access (or multiple access) to the data.
- Impractical to store the whole data.
- Simple calculation per data due to time and space constraints.

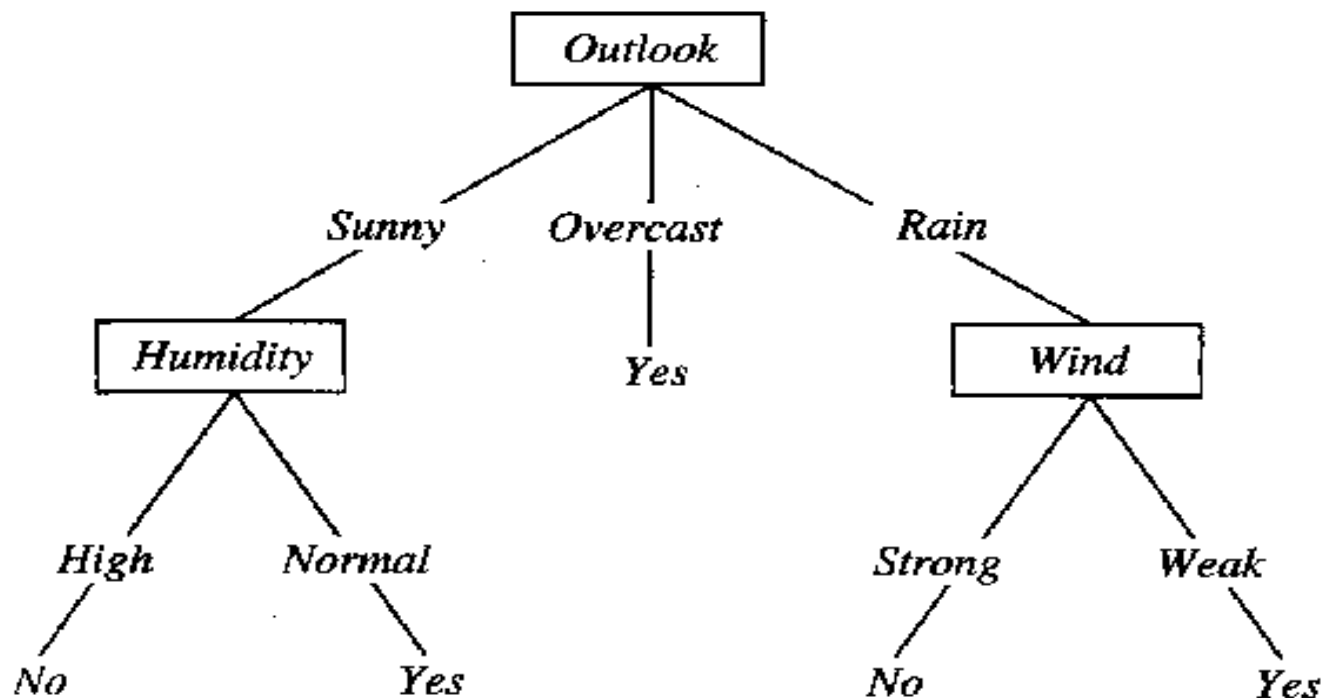
Revision: Tree-Based Classification



- Decision tree is a classification model. Its basic structure is a general tree structure
 - Internal node: test on example's attribute value
 - Leaf node: class labels
- Key idea:
 - 1) pick an attribute to test at root
 - 2) divide the training data into subsets D_i for each value the attribute can take on
 - 3) build the tree for each D_i and splice it in under the appropriate branch at the root

Decision Tree Example

- Shall we go outside today?



Decision Tree – Algorithm



- How to build a decision tree?

Main loop:

1. $A \leftarrow$ the “best” decision attribute for next *node*
2. Assign A as decision attribute for *node*
3. For each value of A , create new descendant of *node*
4. Sort training examples to leaf nodes
5. If training examples perfectly classified, Then STOP, Else iterate over new leaf nodes

Decision Trees – Limitations and Goal



- Limitations
 - Classic decision tree learners assume **all** training data can be **simultaneously stored** in main memory.
 - Disk-based decision tree learners **repeatedly read** training data from disk sequentially
- Goal
 - Design decision tree learners that read each example **at most once**, and use a **small constant time** to process it.

Very Fast Decision Trees (VFDT)



- In order to find the best attribute at a node, it may be sufficient to consider only **a small subset of the training examples** that pass through that node.
 - Given a stream of examples, use the first ones to choose the root attribute.
 - Once the root attribute is chosen, the successive examples are passed down to the corresponding leaves, and used to **choose the attribute** there, and so on recursively.
- Use **Hoeffding bound** to decide how many examples are enough at each node

How to choose an attribute?



- In Session 5, we choose the attributes that can minimise the error rate.
- Recall:
$$\text{Error rate} = \frac{\text{\# wrong predictions}}{\text{total \# of predictions}} = \frac{f_{10} + f_{01}}{f_{11} + f_{10} + f_{01} + f_{00}}$$
- Here we introduce a more algorithmic criterion: **information gain**
- **$\text{Gain}(A) = \text{Info}(D) - \text{Info}_A(D)$**
- $\text{Info}(D)$: the info needed to classify a tuple in D
- $\text{Info}_A(D)$: the info needed to classify D after using A to split D
- $\text{Gain}(A)$: information gain by branching on attribute A
- Calculate the Info Gain for each attribute and pick the one with the largest gain

VFDT (contd.)




- Calculate the **information gain** for the attributes and determines the best two attributes A and B
 - Pre-pruning: consider a “null” attribute that consists of not splitting the node


- At each node, check for the condition

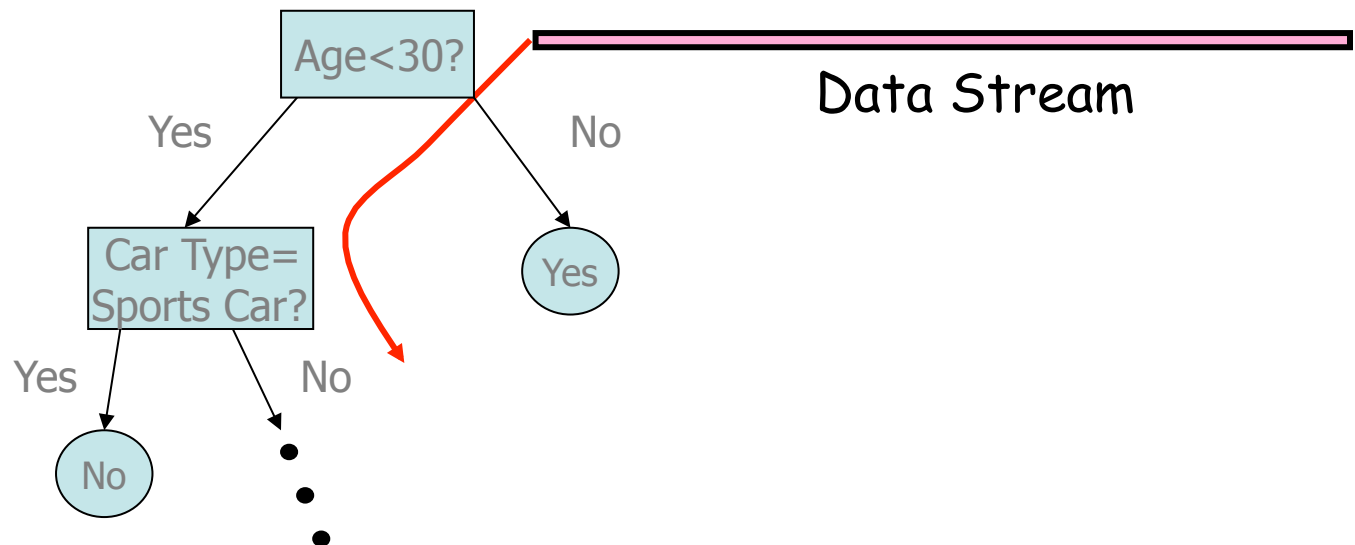
$$\Delta\text{Gain} = \text{Gain}(A) - \text{Gain}(B) > \varepsilon$$

- If the condition is satisfied, create child nodes based on the test at the node
- If not, stream in more examples and perform calculations till condition satisfied

VFDT (contd.)



$$\bar{G}(\text{Car Type}) - \bar{G}(\text{Gender}) > \varepsilon$$




Limitations of VFDT



- VFDT assumes
training data is a sample drawn from *stationary distribution*.
- Most large databases or data streams violate this assumption
 - Concept Drift: data is generated by a *time-changing* concept function, e.g.
 - Seasonal effects
 - Economic cycles
- Goal:
 - Mining continuously changing data streams
 - Scale well

Improvements for VFDT



- Common Approach: Sliding window
 - when a new example arrives, reapply a traditional learner to a sliding window of w most recent examples
- Drawbacks of sliding window
 - Sensitive to window size
 - If w is small relative to the concept shift rate, assure the availability of a model reflecting the current concept
 - Too small w may lead to insufficient examples to learn the concept
 - If examples arrive at a rapid rate or the concept changes quickly, the computational cost of reapplying a learner may be prohibitively high.

CVFDT – Why?



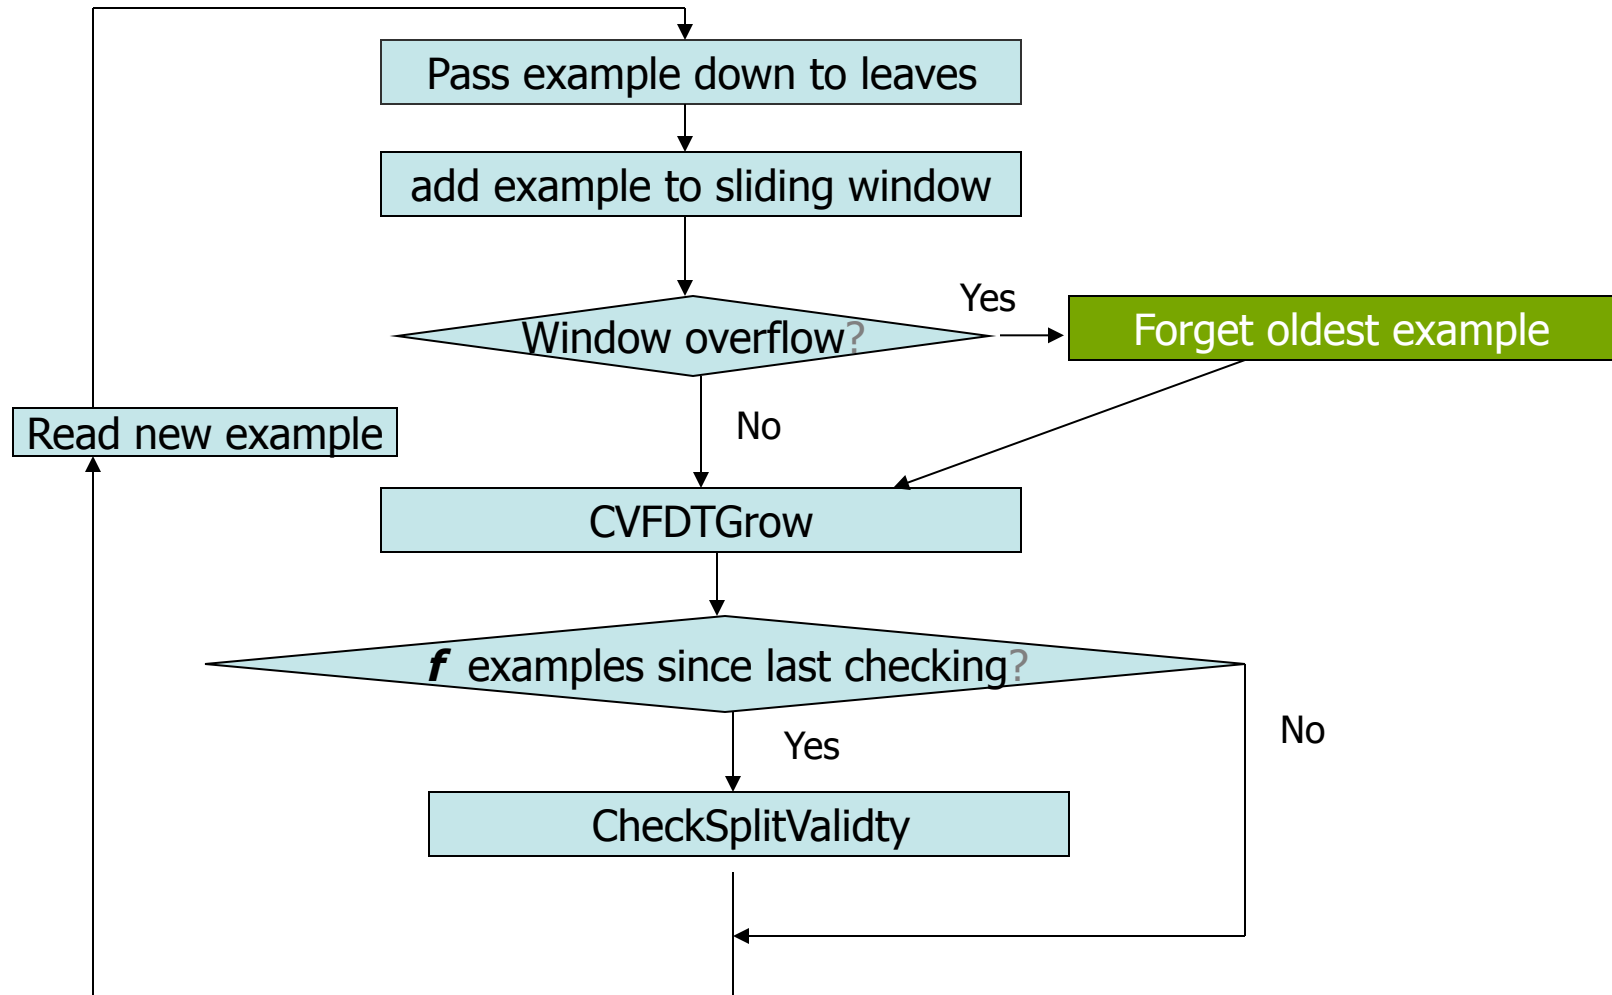
- Concept-adapting very fast decision trees (CVFDT)
 - Extend VFDT
 - Maintain VFDT's speed and accuracy
 - Detect and respond to changes in the example-generating process

CVFDT – How?

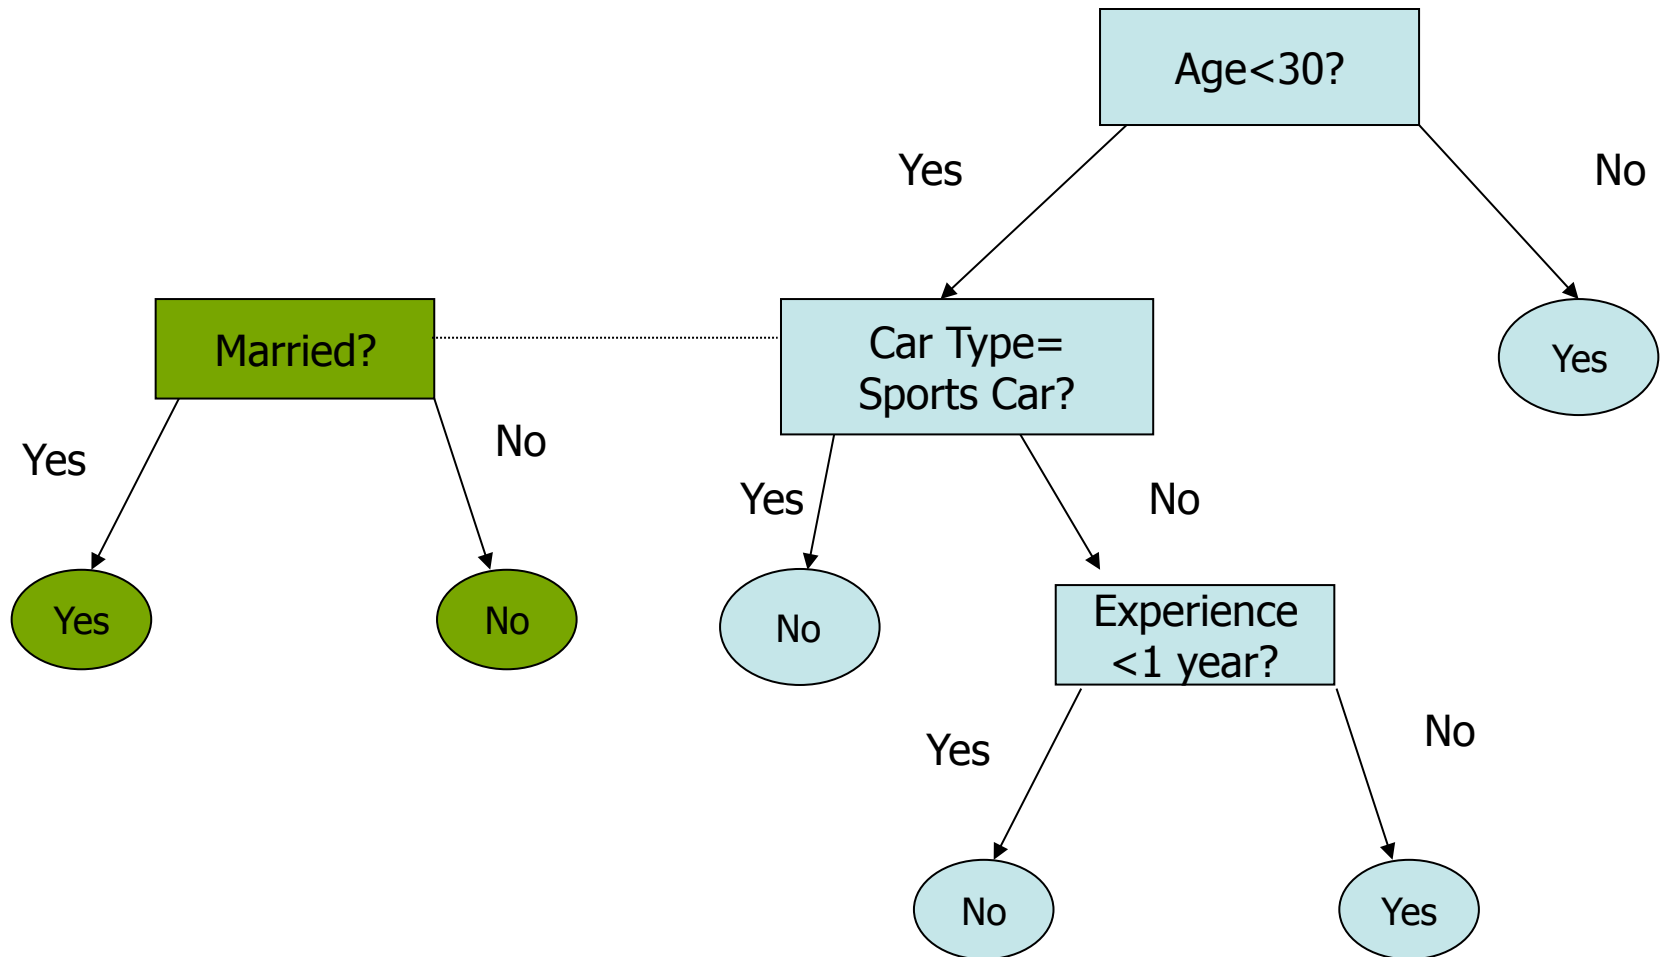


- With a time-changing concept, the current splitting attribute of some nodes may not be the best any more.
- An outdated sub tree may still be better than the best single leaf, particularly if it is near the root.
 - Grow an alternative sub tree with the new best attribute at its root, when the old attribute seems out-of-date.
- Periodically use a bunch of samples to evaluate qualities of trees.
 - Replace the old sub tree when the alternate one becomes more accurate.

CVFDT – The algorithm



CVFDT – Example



CVFDT – Experimental Result

